

dlnd_face_generation

March 27, 2019

1 Face Generation

In this project, you'll define and train a DCGAN on a dataset of faces. Your goal is to get a generator network to generate *new* images of faces that look as realistic as possible!

The project will be broken down into a series of tasks from **loading in data to defining and training adversarial networks**. At the end of the notebook, you'll be able to visualize the results of your trained Generator to see how it performs; your generated samples should look like fairly realistic faces with small amounts of noise.

1.0.1 Get the Data

You'll be using the [CelebFaces Attributes Dataset \(CelebA\)](#) to train your adversarial networks.

This dataset is more complex than the number datasets (like MNIST or SVHN) you've been working with, and so, you should prepare to define deeper networks and train them for a longer time to get good results. It is suggested that you utilize a GPU for training.

1.0.2 Pre-processed Data

Since the project's main focus is on building the GANs, we've done *some* of the pre-processing for you. Each of the CelebA images has been cropped to remove parts of the image that don't include a face, then resized down to 64x64x3 NumPy images. Some sample data is show below.

If you are working locally, you can download this data [by clicking here](#)

This is a zip file that you'll need to extract in the home directory of this notebook for further loading and processing. After extracting the data, you should be left with a directory of data `processed_celeba_small/`

```
In [42]: # can comment out after executing
         #!unzip processed_celeba_small.zip
```

```
In [43]: data_dir = 'processed_celeba_small/'
```

```
"""
DON'T MODIFY ANYTHING IN THIS CELL
"""

import pickle as pkl
import matplotlib.pyplot as plt
```

```
import numpy as np
import problem_unittests as tests
#import helper

%matplotlib inline
```

1.1 Visualize the CelebA Data

The [CelebA](#) dataset contains over 200,000 celebrity images with annotations. Since you're going to be generating faces, you won't need the annotations, you'll only need the images. Note that these are color images with [3 color channels \(RGB\)](#) each.

1.1.1 Pre-process and Load the Data

Since the project's main focus is on building the GANs, we've done *some* of the pre-processing for you. Each of the CelebA images has been cropped to remove parts of the image that don't include a face, then resized down to 64x64x3 NumPy images. This *pre-processed* dataset is a smaller subset of the very large CelebA data.

There are a few other steps that you'll need to **transform** this data and create a **DataLoader**.

Exercise: Complete the following `get_dataloader` function, such that it satisfies these requirements:

- Your images should be square, Tensor images of size `image_size x image_size` in the x and y dimension.
- Your function should return a `DataLoader` that shuffles and batches these Tensor images.

ImageFolder To create a dataset given a directory of images, it's recommended that you use PyTorch's [ImageFolder](#) wrapper, with a root directory `processed_celeba_small/` and data transformation passed in.

```
In [44]: # necessary imports
import torch
from torchvision import datasets
from torchvision import transforms

In [45]: from torch.utils.data import DataLoader

In [46]: import os

In [47]: def get_dataloader(batch_size, image_size, data_dir='processed_celeba_small/'):
    """
    Batch the neural network data using DataLoader
    :param batch_size: The size of each batch; the number of images in a batch
    :param img_size: The square size of the image data (x, y)
    :param data_dir: Directory where image data is located
    :return: DataLoader with batched data
```

```

"""

# define the transformations
transform = transforms.Compose([transforms.Resize(image_size),
                                transforms.ToTensor()])

# create dataset and dataloader
num_workers = 0
train_data = datasets.ImageFolder(root=data_dir, transform=transform)
train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True, num_workers=num_workers)

return train_loader

```

1.2 Create a DataLoader

Exercise: Create a DataLoader celeba_train_loader with appropriate hyperparameters. Call the above function and create a dataloader to view images. * You can decide on any reasonable batch_size parameter * Your image_size **must be** 32. Resizing the data to a smaller size will make for faster training, while still creating convincing images of faces!

```

In [48]: # Define function hyperparameters
batch_size = 64
img_size = 32

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

# Call your function and get a dataloader
celeba_train_loader = get_dataloader(batch_size, img_size)

```

Next, you can view some images! You should see square images of somewhat-centered faces.

Note: You'll need to convert the Tensor images into a NumPy type and transpose the dimensions to correctly display an image, suggested imshow code is below, but it may not be perfect.

```

In [49]: # helper display function
def imshow(img):
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))


"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

# obtain one batch of training images
dataiter = iter(celeba_train_loader)
images, _ = dataiter.next() # _ for no labels

# plot the images in the batch, along with the corresponding labels

```

```
fig = plt.figure(figsize=(20, 4))
plot_size=20
for idx in np.arange(plot_size):
    ax = fig.add_subplot(2, plot_size/2, idx+1, xticks=[], yticks=[])
    imshow(images[idx])
```



Exercise: Pre-process your image data and scale it to a pixel range of -1 to 1 You need to do a bit of pre-processing; you know that the output of a tanh activated generator will contain pixel values in a range from -1 to 1, and so, we need to rescale our training images to a range of -1 to 1. (Right now, they are in a range from 0-1.)

```
In [50]: # TODO: Complete the scale function
def scale(x, feature_range=(-1, 1)):
    ''' Scale takes in an image x and returns that image, scaled
        with a feature_range of pixel values from -1 to 1.
        This function assumes that the input x is already scaled from 0-1. '''
    # assume x is scaled to (0, 1)
    # scale to feature_range and return scaled x

    x_min, x_max = feature_range[0], feature_range[1]
    x = x*(x_max-x_min)-1

    return x
```

```
In [51]: """
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

# check scaled range
# should be close to -1 to 1
img = images[0]
scaled_img = scale(img)

print('Min: ', scaled_img.min())
print('Max: ', scaled_img.max())
```

```
Min: tensor(-0.8824)
Max: tensor(0.8118)
```

2 Define the Model

A GAN is comprised of two adversarial networks, a discriminator and a generator.

2.1 Discriminator

Your first task will be to define the discriminator. This is a convolutional classifier like you've built before, only without any maxpooling layers. To deal with this complex data, it's suggested you use a deep network with **normalization**. You are also allowed to create any helper functions that may be useful.

Exercise: Complete the Discriminator class

- The inputs to the discriminator are 32x32x3 tensor images
- The output should be a single value that will indicate whether a given image is real or fake

```
In [52]: import torch.nn as nn
import torch.nn.functional as F

In [53]: # since we will be using batch normalization layers, a helper function will be useful
def conv(in_channels, out_channels, kernel_size=4, stride=2, padding=1, batch_norm=True):
    """Helper function to create conv layer with opt. batch normalisation. These are
    stride convolutional layers that half input size. Don't need maxpooling layers.
    """

    layers = []

    # need to set the bias to zero when using batch norm layers
    conv_layer = nn.Conv2d(in_channels=in_channels, out_channels=out_channels, kernel_size=kernel_size,
                           stride=stride, padding=padding, bias=False)
    layers.append(conv_layer)

    # optional batch norm layer
    if batch_norm:
        batch_norm = nn.BatchNorm2d(num_features=out_channels)
        layers.append(batch_norm)

    # use the Sequential wrapper
    conv_layers = nn.Sequential(*layers)

    return conv_layers

In [54]: class Discriminator(nn.Module):

    def __init__(self, conv_dim):
        """
        Initialize the Discriminator Module
        :param conv_dim: The depth of the first convolutional layer
        """
```

```

super(Discriminator, self).__init__()

# save class variables
self.conv_dim = conv_dim

# create all layers. Essentially, I just reproduce the architecture from the .
self.conv_1 = conv(3, conv_dim, batch_norm=False) # fist conv layer should ha
# output: 16*16*conv_dim
self.conv_2 = conv(conv_dim, conv_dim*2, batch_norm=True) # output: 8*8*2*con
self.conv_3 = conv(conv_dim*2, conv_dim*4, batch_norm=True) # output: 4*4*4*c
self.fc_1 = nn.Linear(4*4*4*conv_dim, 1, )

def forward(self, x):
    """
    Forward propagation of the neural network
    :param x: The input to the neural network
    :return: Discriminator logits; the output of the neural network
    """

    # convolutional part of the classifier
    x = F.leaky_relu(self.conv_1(x), negative_slope=0.2)
    x = F.leaky_relu(self.conv_2(x), negative_slope=0.2)
    x = F.leaky_relu(self.conv_3(x), negative_slope=0.2)

    # reshape and feed through fully connected layer
    # batch size first
    x = x.view(-1, 4*4*4*self.conv_dim)
    x = self.fc_1(x) # no activation funcito, will use BCE with sigmoid later

    return x

    """
    DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
    """

tests.test_discriminator(Discriminator)

```

Tests Passed

2.2 Generator

The generator should upsample an input and generate a *new* image of the same size as our training data 32x32x3. This should be mostly transpose convolutional layers with normalization applied to the outputs.

Exercise: Complete the Generator class

- The inputs to the generator are vectors of some length `z_size`

- The output should be a image of shape 32x32x3

```
In [55]: # again, useful to create a helper function here
def deconv(in_channels, out_channels, kernel_size=4, stride=2, padding=1, batch_norm=False):
    """This created a transposed convolutional layer with optional batch norm applied
    This will increase input dimensions in every direction by factor 2
    """

    layers = []
    deconv = nn.ConvTranspose2d(in_channels=in_channels, out_channels=out_channels, kernel_size=kernel_size,
                                stride=stride, padding=padding, bias=False)

    layers.append(deconv)

    # optional batch norm layer
    if batch_norm:
        layers.append(nn.BatchNorm2d(num_features=out_channels))

    return nn.Sequential(*layers)

In [56]: class Generator(nn.Module):

    def __init__(self, z_size, conv_dim):
        """
        Initialize the Generator Module
        :param z_size: The length of the input latent vector, z
        :param conv_dim: The depth of the inputs to the *last* transpose convolutional layer
        """
        super(Generator, self).__init__()
        self.conv_dim = conv_dim

        # We start with a fully connected layer to bring z to the right dimensions
        self.fc_1 = nn.Linear(z_size, 4*4*conv_dim*4)

        # couple of deconv layers
        self.deconv_1 = deconv(conv_dim*4, conv_dim*2)
        self.deconv_2 = deconv(conv_dim*2, conv_dim)
        self.deconv_3 = deconv(conv_dim, 3, batch_norm=False)

    def forward(self, x):
        """
        Forward propagation of the neural network
        :param x: The input to the neural network
        :return: A 32x32x3 Tensor image as output
        """
        # feed through fc layers to get to the right size
        x = F.relu(self.fc_1(x))
        x = x.view(-1, self.conv_dim*4, 4, 4)
```

```

        # feed through deconv layers to upsample
        x = F.relu(self.deconv_1(x))
        x = F.relu(self.deconv_2(x))
        x = F.tanh(self.deconv_3(x))

    return x

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""
tests.test_generator(Generator)

```

Tests Passed

2.3 Initialize the weights of your networks

To help your models converge, you should initialize the weights of the convolutional and linear layers in your model. From reading the [original DCGAN paper](#), they say: > All weights were initialized from a zero-centered Normal distribution with standard deviation 0.02.

So, your next task will be to define a weight initialization function that does just this!

You can refer back to the lesson on weight initialization or even consult existing model code, such as that from [the networks.py file in CycleGAN Github repository](#) to help you complete this function.

Exercise: Complete the weight initialization function

- This should initialize only **convolutional** and **linear** layers
- Initialize the weights to a normal distribution, centered around 0, with a standard deviation of 0.02.
- The bias terms, if they exist, may be left alone or set to 0.

In [57]: `from torch.nn import init`

```

In [58]: def weights_init_normal(m):
    """
    Applies initial weights to certain layers in a model .
    The weights are taken from a normal distribution
    with mean = 0, std dev = 0.02.
    :param m: A module or layer in a network
    """

    # classname will be something like:
    # 'Conv', 'BatchNorm2d', 'Linear', etc.
    classname = m.__class__.__name__

    # TODO: Apply initial weights to convolutional and linear layers
    if hasattr(m, 'weight') and (classname.find('Conv') != -1 or classname.find('Line
        init.normal_(m.weight.data, 0.0, 0.02)

```



```

        if hasattr(m, 'bias') and m.bias is not None:
            init.constant_(m.bias.data, 0.0)

```

2.4 Build complete network

Define your models' hyperparameters and instantiate the discriminator and generator from the classes defined above. Make sure you've passed in the correct input arguments.

```

In [59]: """
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

def build_network(d_conv_dim, g_conv_dim, z_size):
    # define discriminator and generator
    D = Discriminator(d_conv_dim)
    G = Generator(z_size=z_size, conv_dim=g_conv_dim)

    # initialize model weights
    D.apply(weights_init_normal)
    G.apply(weights_init_normal)

    print(D)
    print()
    print(G)

    return D, G

```

Exercise: Define model hyperparameters

```

In [60]: # Define model hyperparams
d_conv_dim = 32
g_conv_dim = 32
z_size = 100

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

D, G = build_network(d_conv_dim, g_conv_dim, z_size)

```

```

Discriminator(
  (conv_1): Sequential(
    (0): Conv2d(3, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
  )
  (conv_2): Sequential(
    (0): Conv2d(32, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (conv_3): Sequential(

```

```

        (0): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (fc_1): Linear(in_features=2048, out_features=1, bias=True)
)

Generator(
  (fc_1): Linear(in_features=100, out_features=2048, bias=True)
  (deconv_1): Sequential(
    (0): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (deconv_2): Sequential(
    (0): ConvTranspose2d(64, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (deconv_3): Sequential(
    (0): ConvTranspose2d(32, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
  )
)

```

2.4.1 Training on GPU

Check if you can train on GPU. Here, we'll set this as a boolean variable `train_on_gpu`. Later, you'll be responsible for making sure that `> * Models, * Model inputs, and * Loss function arguments`

Are moved to GPU, where appropriate.

```

In [61]: """
        DON'T MODIFY ANYTHING IN THIS CELL
        """
        import torch

        # Check for a GPU
        train_on_gpu = torch.cuda.is_available()
        if not train_on_gpu:
            print('No GPU found. Please use a GPU to train your neural network.')
        else:
            print('Training on GPU!')

```

Training on GPU!

2.5 Discriminator and Generator Losses

Now we need to calculate the losses for both types of adversarial networks.

2.5.1 Discriminator Losses

- For the discriminator, the total loss is the sum of the losses for real and fake images, $d_loss = d_real_loss + d_fake_loss$.
- Remember that we want the discriminator to output 1 for real images and 0 for fake images, so we need to set up the losses to reflect that.

2.5.2 Generator Loss

The generator loss will look similar only with flipped labels. The generator's goal is to get the discriminator to *think* its generated images are *real*.

Exercise: Complete real and fake loss functions You may choose to use either cross entropy or a least squares error loss to complete the following `real_loss` and `fake_loss` functions.

```
In [62]: train_on_gpu
```

```
Out[62]: True
```

```
In [63]: def real_loss(D_out, smooth=False):
    '''Calculates how close discriminator outputs are to being real.
    param, D_out: discriminator logits
    return: real loss'''

    batch_size = D_out.size(0)
    if smooth:
        labels = torch.ones(batch_size)*0.9
    else:
        labels = torch.ones(batch_size)

    # move to GPU if available
    if train_on_gpu:
        labels = labels.cuda()

    criterion = nn.BCEWithLogitsLoss()
    loss = criterion(D_out.squeeze(), labels)

    return loss

def fake_loss(D_out):
    '''Calculates how close discriminator outputs are to being fake.
    param, D_out: discriminator logits
    return: fake loss'''

    # no label smoothing here
    batch_size = D_out.size(0)
    labels = torch.zeros(batch_size)
    if train_on_gpu:
        labels = labels.cuda()
```

```

criterion = nn.BCEWithLogitsLoss()
loss = criterion (D_out.squeeze(), labels)

return loss

```

2.6 Optimizers

Exercise: Define optimizers for your Discriminator (D) and Generator (G) Define optimizers for your models with appropriate hyperparameters.

```
In [64]: import torch.optim as optim
```

```
In [65]: # Create optimizers for the discriminator D and generator G
```

```

lr = 0.002
beta1 = 0.5
beta2 = 0.999

d_optimizer = optim.Adam(D.parameters(), lr=lr, betas=[beta1, beta2])
g_optimizer = optim.Adam(G.parameters(), lr=lr, betas=[beta1, beta2])

```

2.7 Training

Training will involve alternating between training the discriminator and the generator. You'll use your functions `real_loss` and `fake_loss` to help you calculate the discriminator losses.

- You should train the discriminator by alternating on real and fake images
- Then the generator, which tries to trick the discriminator and should have an opposing loss function

Saving Samples You've been given some code to print out some loss statistics and save some generated "fake" samples.

Exercise: Complete the training function Keep in mind that, if you've moved your models to GPU, you'll also have to move any model inputs to GPU.

```
In [66]: from workspace_utils import active_session
```

```

In [67]: def train(D, G, n_epochs, print_every=50):
        '''Trains adversarial networks for some number of epochs
        param, D: the discriminator network
        param, G: the generator network
        param, n_epochs: number of epochs to train for
        param, print_every: when to print and record the models' losses
        return: D and G losses'''

```

```

# move models to GPU
if train_on_gpu:
    D.cuda()
    G.cuda()

# keep track of loss and generated, "fake" samples
samples = []
losses = []

# Get some fixed data for sampling. These are images that are held
# constant throughout training, and allow us to inspect the model's performance
sample_size=16
fixed_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
fixed_z = torch.from_numpy(fixed_z).float()
# move z to GPU if available
if train_on_gpu:
    fixed_z = fixed_z.cuda()

# epoch training loop
for epoch in range(n_epochs):

    # batch training loop
    for batch_i, (real_images, _) in enumerate(celeba_train_loader):

        batch_size = real_images.size(0)
        real_images = scale(real_images)

        # =====
        #          YOUR CODE HERE: TRAIN THE NETWORKS
        # =====

        # 1. Train the discriminator on real and fake images
        # zero out accumulated gradients
        d_optimizer.zero_grad()

        # move images to the GPU if available
        if train_on_gpu:
            real_images = real_images.cuda()

        # loss on real images
        out_real = D.forward(real_images)
        d_loss_real = real_loss(out_real, smooth=True)

        # loss on fake images
        z = np.random.uniform(-1, 1, size=(batch_size, z_size))
        z = torch.from_numpy(z).float()
        if train_on_gpu:
            z = z.cuda()

```

```

fake_images = G.forward(z)
out_fake = D.forward(fake_images)
d_loss_fake = fake_loss(out_fake)

# sum them up and run a step of optimization
d_loss = d_loss_fake + d_loss_real
d_loss.backward()
d_optimizer.step()

# 2. Train the generator with an adversarial loss
g_optimizer.zero_grad()

# create fake images
z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()
if train_on_gpu:
    z = z.cuda()
fake_images = G.forward(z)
out_fake = D.forward(fake_images)
g_loss = real_loss(out_fake, smooth=True)
g_loss.backward()
g_optimizer.step()

# =====
#                               END OF YOUR CODE
# =====

# Print some loss stats
if batch_i % print_every == 0:
    # append discriminator loss and generator loss
    losses.append((d_loss.item(), g_loss.item()))
    # print discriminator and generator loss
    print('Epoch [{:5d}/{:5d}] | d_loss: {:.4f} | g_loss: {:.4f}'.format(
        epoch+1, n_epochs, d_loss.item(), g_loss.item()))

## AFTER EACH EPOCH##
# this code assumes your generator is named G, feel free to change the name
# generate and save sample, fake images
G.eval() # for generating samples
samples_z = G(fixed_z)
samples.append(samples_z)
G.train() # back to training mode

# Save training generator samples
with open('train_samples.pkl', 'wb') as f:
    pickle.dump(samples, f)

```

```

# save the losses
with open('losses.pkl', 'wb') as f:
    pickle.dump(losses, f)

# finally return losses
return losses

```

Set your number of training epochs and train your GAN!

In [68]: *# define helper functions for saving and loading models*

```

def save_models(filename, D, G):
    filename_G = os.path.splitext(os.path.basename(filename))[0] + '_G.pt'
    filename_D = os.path.splitext(os.path.basename(filename))[0] + '_D.pt'

    torch.save(D, os.path.join('saved_models', filename_G))
    torch.save(G, os.path.join('saved_models', filename_D))

```

In [69]: *def load_models(filename):*

```

    filename_G = os.path.splitext(os.path.basename(filename))[0] + '_G.pt'
    filename_D = os.path.splitext(os.path.basename(filename))[0] + '_D.pt'

    D = torch.load(os.path.join('saved_models', filename_D))
    G = torch.load(os.path.join('saved_models', filename_G))

    return D, G

```

In [70]: *# set number of epochs*

```
n_epochs = 20
```

```

"""
DON'T MODIFY ANYTHING IN THIS CELL
"""

# call training function
with active_session():
    losses = train(D, G, n_epochs=n_epochs, print_every=200)
    # save models
    save_models('test_run', D, G)

```

```

Epoch [ 1/ 20] | d_loss: 1.4085 | g_loss: 1.6797
Epoch [ 1/ 20] | d_loss: 1.9654 | g_loss: 2.8596
Epoch [ 1/ 20] | d_loss: 0.9133 | g_loss: 3.2011
Epoch [ 1/ 20] | d_loss: 1.2665 | g_loss: 2.3290
Epoch [ 1/ 20] | d_loss: 0.8755 | g_loss: 2.9497
Epoch [ 1/ 20] | d_loss: 0.9073 | g_loss: 2.1558
Epoch [ 1/ 20] | d_loss: 1.0728 | g_loss: 1.8908
Epoch [ 1/ 20] | d_loss: 1.1615 | g_loss: 1.7977

```

Epoch [2/	20]	d_loss: 1.0007	g_loss: 2.0010
Epoch [2/	20]	d_loss: 1.2092	g_loss: 1.0342
Epoch [2/	20]	d_loss: 1.1393	g_loss: 2.9866
Epoch [2/	20]	d_loss: 1.2985	g_loss: 2.7369
Epoch [2/	20]	d_loss: 0.9051	g_loss: 1.6610
Epoch [2/	20]	d_loss: 1.0236	g_loss: 1.9128
Epoch [2/	20]	d_loss: 1.1514	g_loss: 1.4379
Epoch [2/	20]	d_loss: 1.0792	g_loss: 1.1933
Epoch [3/	20]	d_loss: 1.1338	g_loss: 1.0980
Epoch [3/	20]	d_loss: 1.1549	g_loss: 2.1727
Epoch [3/	20]	d_loss: 1.1001	g_loss: 1.4267
Epoch [3/	20]	d_loss: 1.0624	g_loss: 1.8107
Epoch [3/	20]	d_loss: 1.2498	g_loss: 0.9424
Epoch [3/	20]	d_loss: 1.1725	g_loss: 1.0707
Epoch [3/	20]	d_loss: 1.1299	g_loss: 2.1788
Epoch [3/	20]	d_loss: 1.2431	g_loss: 0.9933
Epoch [4/	20]	d_loss: 1.2462	g_loss: 1.8119
Epoch [4/	20]	d_loss: 1.1436	g_loss: 1.4189
Epoch [4/	20]	d_loss: 1.1219	g_loss: 1.7847
Epoch [4/	20]	d_loss: 1.2202	g_loss: 0.9966
Epoch [4/	20]	d_loss: 1.1492	g_loss: 0.9284
Epoch [4/	20]	d_loss: 1.1970	g_loss: 1.4127
Epoch [4/	20]	d_loss: 1.2645	g_loss: 1.0563
Epoch [4/	20]	d_loss: 1.0906	g_loss: 1.4178
Epoch [5/	20]	d_loss: 1.0363	g_loss: 1.1238
Epoch [5/	20]	d_loss: 1.1172	g_loss: 1.2252
Epoch [5/	20]	d_loss: 1.0411	g_loss: 1.4810
Epoch [5/	20]	d_loss: 1.1754	g_loss: 1.4299
Epoch [5/	20]	d_loss: 1.0971	g_loss: 1.5812
Epoch [5/	20]	d_loss: 1.2562	g_loss: 1.3027
Epoch [5/	20]	d_loss: 1.1893	g_loss: 1.5567
Epoch [5/	20]	d_loss: 1.0495	g_loss: 0.9759
Epoch [6/	20]	d_loss: 1.0843	g_loss: 1.4443
Epoch [6/	20]	d_loss: 0.9862	g_loss: 1.2099
Epoch [6/	20]	d_loss: 1.3647	g_loss: 1.9470
Epoch [6/	20]	d_loss: 1.3128	g_loss: 2.0859
Epoch [6/	20]	d_loss: 0.9008	g_loss: 1.7177
Epoch [6/	20]	d_loss: 0.9097	g_loss: 1.8823
Epoch [6/	20]	d_loss: 1.0772	g_loss: 1.8933
Epoch [6/	20]	d_loss: 0.9980	g_loss: 1.2596
Epoch [7/	20]	d_loss: 1.1620	g_loss: 1.6711
Epoch [7/	20]	d_loss: 0.9617	g_loss: 1.8993
Epoch [7/	20]	d_loss: 0.9540	g_loss: 1.3879
Epoch [7/	20]	d_loss: 0.8085	g_loss: 2.0945
Epoch [7/	20]	d_loss: 1.0186	g_loss: 1.5824
Epoch [7/	20]	d_loss: 1.0008	g_loss: 2.2853
Epoch [7/	20]	d_loss: 1.0539	g_loss: 2.1235
Epoch [7/	20]	d_loss: 0.9257	g_loss: 1.3333

Epoch [8/	20]	d_loss: 1.3719	g_loss: 2.5110
Epoch [8/	20]	d_loss: 1.0630	g_loss: 2.1658
Epoch [8/	20]	d_loss: 0.8910	g_loss: 1.3444
Epoch [8/	20]	d_loss: 1.1450	g_loss: 1.5528
Epoch [8/	20]	d_loss: 0.9285	g_loss: 1.3445
Epoch [8/	20]	d_loss: 0.8835	g_loss: 1.8925
Epoch [8/	20]	d_loss: 1.0031	g_loss: 1.2902
Epoch [8/	20]	d_loss: 1.1903	g_loss: 1.5033
Epoch [9/	20]	d_loss: 0.9921	g_loss: 1.2919
Epoch [9/	20]	d_loss: 1.7581	g_loss: 2.3276
Epoch [9/	20]	d_loss: 1.1013	g_loss: 1.2751
Epoch [9/	20]	d_loss: 0.7197	g_loss: 2.6704
Epoch [9/	20]	d_loss: 1.0329	g_loss: 2.3747
Epoch [9/	20]	d_loss: 0.9465	g_loss: 2.0504
Epoch [9/	20]	d_loss: 0.9179	g_loss: 1.7869
Epoch [9/	20]	d_loss: 0.8754	g_loss: 2.5746
Epoch [10/	20]	d_loss: 1.0108	g_loss: 2.0277
Epoch [10/	20]	d_loss: 1.0832	g_loss: 2.0219
Epoch [10/	20]	d_loss: 1.0131	g_loss: 2.1754
Epoch [10/	20]	d_loss: 0.9172	g_loss: 2.0244
Epoch [10/	20]	d_loss: 1.0627	g_loss: 1.9043
Epoch [10/	20]	d_loss: 0.9397	g_loss: 2.2459
Epoch [10/	20]	d_loss: 0.8826	g_loss: 2.1297
Epoch [10/	20]	d_loss: 0.7076	g_loss: 2.3231
Epoch [11/	20]	d_loss: 1.4315	g_loss: 2.8529
Epoch [11/	20]	d_loss: 0.8496	g_loss: 2.0763
Epoch [11/	20]	d_loss: 0.7054	g_loss: 2.8200
Epoch [11/	20]	d_loss: 0.9234	g_loss: 1.5576
Epoch [11/	20]	d_loss: 0.8166	g_loss: 1.5277
Epoch [11/	20]	d_loss: 0.7818	g_loss: 1.7172
Epoch [11/	20]	d_loss: 1.1832	g_loss: 2.0288
Epoch [11/	20]	d_loss: 0.9384	g_loss: 1.4757
Epoch [12/	20]	d_loss: 1.0270	g_loss: 1.3097
Epoch [12/	20]	d_loss: 0.8983	g_loss: 2.0235
Epoch [12/	20]	d_loss: 0.9679	g_loss: 2.6909
Epoch [12/	20]	d_loss: 0.9389	g_loss: 1.1859
Epoch [12/	20]	d_loss: 0.9258	g_loss: 1.4827
Epoch [12/	20]	d_loss: 0.9246	g_loss: 0.9269
Epoch [12/	20]	d_loss: 0.8519	g_loss: 1.7244
Epoch [12/	20]	d_loss: 0.7193	g_loss: 1.4648
Epoch [13/	20]	d_loss: 1.6197	g_loss: 1.1044
Epoch [13/	20]	d_loss: 0.9497	g_loss: 1.8480
Epoch [13/	20]	d_loss: 0.5861	g_loss: 2.3280
Epoch [13/	20]	d_loss: 0.9832	g_loss: 2.0721
Epoch [13/	20]	d_loss: 0.7572	g_loss: 1.6757
Epoch [13/	20]	d_loss: 1.0164	g_loss: 1.8160
Epoch [13/	20]	d_loss: 0.8732	g_loss: 1.0368
Epoch [13/	20]	d_loss: 0.9519	g_loss: 2.1948

Epoch [14/	20]	d_loss: 1.0249	g_loss: 2.0040
Epoch [14/	20]	d_loss: 0.7439	g_loss: 1.7291
Epoch [14/	20]	d_loss: 0.5619	g_loss: 2.3486
Epoch [14/	20]	d_loss: 0.9228	g_loss: 1.3780
Epoch [14/	20]	d_loss: 1.3818	g_loss: 0.4998
Epoch [14/	20]	d_loss: 0.9649	g_loss: 1.9225
Epoch [14/	20]	d_loss: 0.9694	g_loss: 1.9332
Epoch [14/	20]	d_loss: 0.9179	g_loss: 2.4936
Epoch [15/	20]	d_loss: 1.4050	g_loss: 4.0071
Epoch [15/	20]	d_loss: 0.7551	g_loss: 1.8494
Epoch [15/	20]	d_loss: 0.8381	g_loss: 1.3455
Epoch [15/	20]	d_loss: 0.6899	g_loss: 2.2511
Epoch [15/	20]	d_loss: 0.7439	g_loss: 2.0075
Epoch [15/	20]	d_loss: 1.0011	g_loss: 1.8389
Epoch [15/	20]	d_loss: 0.6235	g_loss: 1.8882
Epoch [15/	20]	d_loss: 0.9006	g_loss: 1.3794
Epoch [16/	20]	d_loss: 0.6907	g_loss: 2.2416
Epoch [16/	20]	d_loss: 1.3841	g_loss: 1.1147
Epoch [16/	20]	d_loss: 0.7801	g_loss: 2.2977
Epoch [16/	20]	d_loss: 0.6376	g_loss: 2.6631
Epoch [16/	20]	d_loss: 0.6358	g_loss: 2.7277
Epoch [16/	20]	d_loss: 0.9399	g_loss: 1.7749
Epoch [16/	20]	d_loss: 0.7202	g_loss: 1.7689
Epoch [16/	20]	d_loss: 1.0842	g_loss: 2.9711
Epoch [17/	20]	d_loss: 0.9619	g_loss: 2.4094
Epoch [17/	20]	d_loss: 0.6688	g_loss: 2.6273
Epoch [17/	20]	d_loss: 1.4767	g_loss: 3.8651
Epoch [17/	20]	d_loss: 0.8388	g_loss: 2.5390
Epoch [17/	20]	d_loss: 1.5722	g_loss: 1.1288
Epoch [17/	20]	d_loss: 0.7022	g_loss: 2.7000
Epoch [17/	20]	d_loss: 2.5118	g_loss: 5.6599
Epoch [17/	20]	d_loss: 0.6498	g_loss: 2.6023
Epoch [18/	20]	d_loss: 1.1505	g_loss: 1.6230
Epoch [18/	20]	d_loss: 0.6662	g_loss: 1.4589
Epoch [18/	20]	d_loss: 0.6843	g_loss: 1.8510
Epoch [18/	20]	d_loss: 0.7024	g_loss: 2.4766
Epoch [18/	20]	d_loss: 0.7368	g_loss: 2.4390
Epoch [18/	20]	d_loss: 0.6951	g_loss: 2.1719
Epoch [18/	20]	d_loss: 0.8790	g_loss: 2.0607
Epoch [18/	20]	d_loss: 0.7513	g_loss: 2.8165
Epoch [19/	20]	d_loss: 0.8602	g_loss: 1.9574
Epoch [19/	20]	d_loss: 0.6102	g_loss: 2.9343
Epoch [19/	20]	d_loss: 0.8086	g_loss: 1.6789
Epoch [19/	20]	d_loss: 0.5976	g_loss: 2.0870
Epoch [19/	20]	d_loss: 1.4322	g_loss: 2.0248
Epoch [19/	20]	d_loss: 0.7777	g_loss: 2.0528
Epoch [19/	20]	d_loss: 0.6091	g_loss: 3.1853
Epoch [19/	20]	d_loss: 0.7229	g_loss: 1.8909

```
Epoch [ 20/ 20] | d_loss: 1.3079 | g_loss: 3.7941
Epoch [ 20/ 20] | d_loss: 0.6446 | g_loss: 2.0223
Epoch [ 20/ 20] | d_loss: 0.7095 | g_loss: 2.0462
Epoch [ 20/ 20] | d_loss: 0.8128 | g_loss: 2.8174
Epoch [ 20/ 20] | d_loss: 0.9917 | g_loss: 2.4630
Epoch [ 20/ 20] | d_loss: 0.9123 | g_loss: 1.7819
Epoch [ 20/ 20] | d_loss: 0.7884 | g_loss: 3.0656
Epoch [ 20/ 20] | d_loss: 1.3286 | g_loss: 0.8428
```

```
/opt/conda/lib/python3.6/site-packages/torch/serialization.py:193: UserWarning: Couldn't retrieve
"__type__" + obj.__name__ + ". It won't be checked "
/opt/conda/lib/python3.6/site-packages/torch/serialization.py:193: UserWarning: Couldn't retrieve
"__type__" + obj.__name__ + ". It won't be checked "
```


2.8 Training loss

Plot the training losses for the generator and discriminator, recorded after each epoch.

```
In [71]: # load in the models
         # D, G = load_models('test_save')

In [72]: fig, ax = plt.subplots()
         losses = np.array(losses)
         plt.plot(losses.T[0], label='Discriminator', alpha=0.5)
         plt.plot(losses.T[1], label='Generator', alpha=0.5)
         plt.title("Training Losses")
         plt.legend()

Out[72]: <matplotlib.legend.Legend at 0x7f135ad78f60>
```



```
In [73]: #losses_no_smoothing = losses
```

```
In [74]: #losses_smoothing = losses
```

2.9 Generator samples from training

View samples of images from the generator, and answer a question about the strengths and weaknesses of your trained models.


```
In [75]: # helper function for viewing a list of passed in sample images
```

```
def view_samples(epoch, samples):
    fig, axes = plt.subplots(figsize=(16,4), nrows=2, ncols=8, sharey=True, sharex=True)
    for ax, img in zip(axes.flatten(), samples[epoch]):
        img = img.detach().cpu().numpy()
        img = np.transpose(img, (1, 2, 0))
        img = ((img + 1)*255 / (2)).astype(np.uint8)
        ax.xaxis.set_visible(False)
        ax.yaxis.set_visible(False)
        im = ax.imshow(img.reshape((32,32,3)))
```

```
In [76]: # Load samples from generator, taken while training
```

```
with open('train_samples.pkl', 'rb') as f:
    samples = pkl.load(f)
```

```
In [77]: _ = view_samples(-1, samples)
```



2.9.1 Question: What do you notice about your generated samples and how might you improve this model?

When you answer this question, consider the following factors: * The dataset is biased; it is made of "celebrity" faces that are mostly white * Model size; larger models have the opportunity to learn more features in a data feature space * Optimization strategy; optimizers and number of epochs affect your final result

Answer:

What do I notice about the generated samples? * first of all, they do look like faces (mostly) which is good * a bit blurry * sometimes looks like two faces merged into one

How might I improve this model? * Larger, more balanced dataset would allow one to create more realistic and more diverse faces * Increasing the model size: more convolutional layers! Potentially using a pre-trained classifier. Would however make it difficult for the generator to keep up in the learning process * Try different loss functions, mean squared loss for example, as we saw in the lectures * Run for more epochs, potentially try SGD or other optimizers, play around with the model parameters

2.9.2 Submitting This Project

When submitting this project, make sure to run all the cells before saving the notebook. Save the notebook file as "dlnd_face_generation.ipynb" and save it as a HTML file under "File" -> "Download as". Include the "problem_unittests.py" files in your submission.